

Neural Reservoir Surrogate on the Norne Field — X-MeshGraphNet

Author: Dr. Montaser Ramadan · **Date:** 2026-06-10 · **Stack:** NVIDIA PhysicsNeMo 1.3.0 / X-MeshGraphNet

Summary. I trained an X-MeshGraphNet (X-MGN) graph-network surrogate to roll the Norne reservoir forward one report-step at a time, predicting **PRESSURE** and **SWAT** on the native 44,431-active-cell Eclipse grid. On the 6 held-out (seed-42) test cases the surrogate reaches a **single-step** validation accuracy of **0.89 bar** PRESSURE / **0.0029** SWAT RMSE, and a **fully autoregressive** 62-step rollout accuracy of **9.27 bar** PRESSURE / **0.0128** SWAT RMSE ($\approx 3.3\%$ of the ~ 278 bar field). The main finding is the $\approx 10\times$ **gap between those two numbers**: it is *autoregressive drift*, it is **non-monotonic** (it peaks mid-rollout around the most active water-front period, not at the final step), and spatially it concentrates at the **advancing SWAT front in the NE fault compartment**, which is also the region of highest observed prediction error.

1. Setup & environment notes

- **Hardware:** 1× NVIDIA RTX 4090 (24 GB), 32 vCPU, 124 GB RAM, RunPod, Ubuntu 24.04, driver 580.
- **Storage:** the dataset and all outputs live on a **native ext4 200 GB volume** mounted at `/workspace`. This was deliberate: the preprocessor materialises ~ 49 GB of partition tensors, and a 9p/network-FUSE path starves the GPU dataloader. The 124 GB of RAM page-caches the partition set after the first epoch, so the network volume's small-file latency stops mattering.
- **Stack:** Python 3.10, `torch 2.4.0+cu121`, `nvidia-physicsnemo 1.3.0`, PyG (cu121), `resdata 6.0.1` for Eclipse I/O.
- **Issues hit & resolved (abbreviated; full log in `TROUBLESHOOTING.md` / `PATCHES.md`):**
 - *MLflow file-store block.* MLflow ≥ 3 refuses a bare file store; set `MLFLOW_ALLOW_FILE_STORE=true` before `train.py` / `inference.py` (otherwise it raises at init).
 - *Dataloader starvation.* `num_workers` was hard-coded to 0; with the partition set on a network volume the GPU sat at 45 %. Setting `num_workers=8`, `persistent_workers=True`, `prefetch_factor=4` raised GPU utilisation to $\sim 96\%$ and cut epoch time **13.7** → **3.5 min ($\approx 4\times$)**, result-neutral.
 - *Inference reordering bug (corrected; see §6.1).* The natural-order reorder was applied per *partition* rather than to the concatenated full-case array, raising a `(14941, 2)` vs `(44431, 2)` broadcast error. Corrected by concatenating the partition outputs before applying the permutation.
- **Patches relied on:**
 - **Patch #5 — case-name parsing.** The stock parser only handled 4-part filenames (`CASE_2D_1_000`); Norne's 3-part names (`NORNE_002_012`) fell through to a fallback that treated every `(case, timestep)` as its own single-timestep "case" — silently **disabling autoregressive rollout**. The 3-part branch restores it. The inference log line `case NORNE_002 (62 timesteps)` is the verification that rollout is active.

- **Patch #7 — cell ordering.** Predictions come out in METIS-partition order; without reordering to EGRID natural (active-cell) order, every per-cell value lands at the wrong physical location downstream (ResInsight, UNRST). Each output HDF5 now carries `attrs["ordering"]=="natural"`, and the spatial figure (§6) is the visual proof the reorder is correct.

2. Data & the mandated split

- **60 Norne LHS cases**, grid **46×112×22 (44,431 active cells)**, **65 report steps** (~9 years). The uncertainty axis is a Latin-hypercube sweep of **fault transmissibility multipliers**: the cases differ in how sealing/conductive the faults are, which motivates a fault-aware graph as the inductive bias (see §7).
- **Split:** seeded 80/10/10 (`random_seed: 42`), **by case** (no timestep leakage). The assignment produced by our run (read back from the partition manifests) matches the mandated seed-42 split exactly:
 - **Test (6):** NORNE_002, 008, 016, 018, 041, 048
 - **Val (6):** NORNE_007, 009, 015, 044, 052, 057
 - **Train (48):** the remaining cases (001, 003–006, 010–014, 017, 019–040, 042–043, 045–047, 049–051, 053–056, 058–060).
- Graph build: **3,720 graphs** (60 cases × 62 usable steps, since `prev_timesteps=2` consumes the first two steps as history). The preprocessor's silent-skip behaviours (cases with <3 steps dropped; `log10` floored at 1e-10; METIS→SimplePartition fallback) were all explicitly gated: graph count, per-split counts, 0/NaN check on a partition sample, and the METIS confirmation in the prep log.

3. Model & configuration

- **X-MGN:** 5 message-passing layers, hidden 128, SiLU, gradient-checkpointed (`checkpoint_segments=2`).
- **Node features (~17–19 d):** static `PERMX`, `PORV`, `X`, `Y`, `Z`; dynamic `PRESSURE`, `SWAT`, `WCID` at the current + 2 previous steps; globals `delta_t`, `time`.
- **Edge features:** `TRANX`, `TRANY`, `TRANZ`, `TRANNNC` (`log10`). `TRANNNC` carries the **non-neighbour connections (NNCs)** that encode the faults, which is the primary reason X-MGN suits Norne.
- **Targets:** `PRESSURE`, `SWAT`, weights `[1, 1]`, losses `[L2, L1]` on *normalised* values; `PERMX` / `TRAN` `log10`-scaled (permeability spans ~6 orders of magnitude).
- **Optimiser:** AdamW, `lr 1e-3 → 1e-6`, weight decay 1e-3, **CosineAnnealingLR** with `T_max=num_epochs`.

Config changes from the canonical Norne baseline, with reasoning:

Change	From → To	Why
<code>num_epochs</code>	1000 → 150	The LR schedule is <code>CosineAnnealingLR(T_max=num_epochs)</code> . At 1000, a realistic ~150-epoch run leaves the LR at ~99 % of its start — the schedule never decays . Setting <code>T_max=150</code> lets the cosine fully anneal within the compute budget (~9 h on one 4090) and is ample for 48-case convergence.
<code>early_stopping.patience</code>	20 → 40	Prevent early stopping from cutting the cosine decay short before the low-LR refinement phase.
<code>simulator</code>	OPM → ECLIPSE	The supplied bundle was run through Eclipse (per <code>data/README</code>), so the readers must use the Eclipse keyword conventions.
<code>random_seed</code>	implicit → explicit 42	Mandated and reproducibility — pins the by-case split.

4. Training run

- **150 epochs, ~3.5 min/epoch, ≈9 h 08 m wall-clock** (15:05 → 00:13), batch size 1 (one graph-partition at a time — a GNN-memory constraint), 2,976 train / 372 val samples.
- **Convergence:** the single-step validation loss falls smoothly and monotonically; the cosine LR reaches ~1e-6 by epoch 150 and the curve has clearly flattened. **Best val loss 2.065e-3** at epoch 150 (early stopping never triggered — the run used its full budget by design).
- **Single-step validation curve (denormalised, teacher-forced):**

Epoch	val_loss	PRESSURE RMSE (bar)	SWAT RMSE
5	3.59e-2	7.68	0.0115
25	1.28e-2	3.60	0.0079
55	6.54e-3	2.17	0.0050
95	3.78e-3	1.16	0.0035
125	2.63e-3	0.914	0.0030
150	2.07e-3	0.889	0.0029

(These are single-step metrics — the model sees the true previous two steps. The autoregressive test numbers in §5 are substantially larger; that gap is the dominant error source.)

5. Held-out (test) performance

Two regimes, and the distinction matters:

- **Single-step (teacher-forced):** PRESSURE **0.889 bar**, SWAT **0.0029** RMSE (epoch-150 validation above) — the model is an excellent one-step operator.
- **Fully autoregressive (62-step rollout; 3 true steps, then the model’s own output is fed back):** aggregated over the 6 test cases, in **physical units**:

Variable	RMSE (norm)	RMSE (phys)	MAE (phys)	worst case
PRESSURE	1.41e-2	9.27 bar ($\approx 3.3\%$ of ~ 278 bar)	4.64 bar	9.86 bar (NORNE_002)
SWAT	8.1e-3	0.0128	0.0030	0.0132 (NORNE_016)

- **Autoregressive drift (error vs rollout step):** *non-monotonic*. PRESSURE RMSE climbs from 0.27 bar at step 3 to a **peak ~ 12.2 bar at step 13**, recovers to ~ 3.8 bar by step 23, then drifts back up to ~ 10.5 bar by step 63. SWAT mirrors this, peaking (~ 0.036) at step 13. The peak is not at the end of the rollout — it coincides with the **most dynamic flow period** (early water-front movement / well rate changes), when one-step errors are largest and the feedback loop amplifies them fastest. Once the field stabilises, the model partially re-tracks the truth.


The $\sim 10\times$ single-step→autoregressive gap is the price of closed-loop rollout: small per-step biases compound. This is the dominant error source — **not** raw model capacity (the single-step operator is already at sub-bar accuracy).

6. Failure Mode Analysis

6.1 Inference Partition-Reordering Bug

The first inference run crashed: `_reorder_to_natural1` did `out[perm] = arr` with `perm` a **full-grid** permutation (44,431) but `arr` a **single partition** ($\sim 14,941$). The reorder was being applied to each partition array individually, when `perm` (built by concatenating every partition’s inner-node global indices) indexes the **concatenated** full-case array. Fix: concatenate the partition outputs first, then reorder. This indicates that the inference loop emits values *per partition in METIS order*, and the correct point to map back to EGRID-natural order is after the partitions are joined. Left unfixed, this would not have raised an error in a different code path; it would have produced spatially scrambled UNRSTs with *correct aggregate metrics* — a bug that preserves aggregate metrics while corrupting spatial outputs, which is why Patch #7 writes an explicit `ordering` attribute as a tripwire.

6.2 Where the surrogate fails, spatially



TRUE/PRED/DIFF panel — final-timestep top-down maps, best→median→worst test cases.

Columns: PRESSURE truth/pred/diff, SWAT truth/pred/diff.

The TRUE/PRED/DIFF panel (`inference/final_timestep_matrix.png` , best→median→worst rows) shows: - **PRESSURE** truth and prediction are visually almost indistinguishable; the diff column is near-zero over the bulk of the field with faint structure near the high-rate wells. - **SWAT** error is **spatially concentrated at the advancing water front in the NE compartment** — a coherent red/blue band that recurs across all 6 test cases, not random speckle. The worst case (NORNE_002) has the largest front displacement error.

Hypothesis. Sharp saturation fronts are what message-passing tends to smooth: a 5-hop receptive field averages over the few cells where SWAT jumps from ~ 0.2 to ~ 0.8 , so the predicted front is slightly diffuse and slightly mis-timed. Under autoregression that small front-position error feeds back and compounds, which is also why the drift peaks during the active-flooding window (§5) rather than at the end. The faults modulate where the front can go, so the error organises along fault-bounded compartments. **Supporting quantitative cut:** the error is carried by SWAT, not PRESSURE — PRESSURE is a smooth, globally-coupled (elliptic-like) field the network reproduces to sub-bar accuracy, whereas SWAT is a sharp, advective front; the per-variable split (PRESSURE 3.3 % vs SWAT front-localised) isolates the front as the failure locus.

6.3 Expected Impact of a Loss-Weighting Change

The diagnosis in §6.2 points to a specific, testable fix. The loss weights PRESSURE-L2 and SWAT-L1 **equally** [1,1] **on normalised values**, but PRESSURE is the smooth, globally-coupled field while SWAT carries most of the front error, so on a shared-gradient multi-task objective the pressure objective may dominate the optimization and the front becomes under-weighted. **Proposed change: up-weight SWAT, e.g.** [1,1] $\rightarrow$ [1,3], holding everything else fixed (split, normalisation, seed, schedule). **Expected outcome:** SWAT front RMSE drops while PRESSURE RMSE rises slightly — the standard multi-task trade-off — which would confirm the front was under-weighted. If instead PRESSURE regressed with SWAT flat, the cause would be the 5-hop receptive field / front sharpness rather than the loss weighting, and the next step would be a multi-scale graph or a gradient-aware loss (§8.2).

7. FNO vs. X-MGN for Faulted Reservoirs

Norne's defining feature is its **sealing/partially-sealing faults**, represented as **NNCs** — connections between cells that are *not* spatial neighbours, and *non*-connections between cells that *are*. A **Fourier Neural Operator** assumes a regular grid and learns smooth, global spectral kernels; it has no way to represent "these two adjacent cells are hydraulically disconnected." It will therefore **blur flow across sealing faults**, leaking pressure and saturation through barriers that the physics keeps shut — most severely in the NE compartment where the SWAT error already concentrates. **X-MGN** encodes each connection as a graph edge weighted by transmissibility (incl. `TRANNNC`), so the fault topology is *in the model*, not approximated away. Where FNO wins: smooth, structured, unfaulted fields and raw throughput (spectral convs are cheap and the spectral bias suits elliptic PDEs like single-phase pressure). It is weaker on sharp fronts across irregular fault boundaries, which is the regime Norne represents.

8. Future Work

1. **Curb autoregressive drift directly:** pushforward / noise-injection training (feed the model its own noised predictions during training) so it learns to correct its own error — the standard fix for the §5 compounding, and likely higher-leverage than any architecture change here.
2. **Sharpen the front:** a gradient-aware or front-weighted loss term, or a multi-scale graph so SWAT shocks aren't averaged across a 5-hop neighbourhood.
3. **Physics guards:** soft mass-balance and a hard `SWAT $\in$ [0,1-SGAS]` constraint (already applied at the UNRST-export stage; better learned in-loop).
4. **More cases / UQ:** the 48-case train set is small for the fault-multiplier space; more LHS samples plus an ensemble would give calibrated uncertainty across the fault-sealing axis.
5. **3-phase:** add SGAS as a third target for full saturation-state prediction.

All numbers trace to committed artifacts: `inference/NORNE_*.hdf5`, `inference/per_case_metrics.csv`, `inference/error_vs_timestep.csv`, `inference/summary_matrix.txt`, `inference/final_timestep_matrix.png`, the run-1 training log, and the seed-42 partition manifests.